

```
In [33]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input (

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when y
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current sess.
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]


```
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00025.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00082.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00001.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00043.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00005.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00048.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00007.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00098.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00089.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00038.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00087.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00085.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00075.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00064.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00068.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00041.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00050.wav  
/kaggle/input/musicgenreclassification/genres_original/pop/pop.00032.wav
```

In [34]: *# Music Genre Classification using Raw Waveform and Pretrained Models*

For Kaggle environment

```
import os  
import numpy as np  
import pandas as pd  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import Dataset, DataLoader  
import torchaudio  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import confusion_matrix, classification_report  
from tqdm.notebook import tqdm  
from transformers import Wav2Vec2Model, Wav2Vec2Config  
  
# Set seeds for reproducibility  
torch.manual_seed(42)  
np.random.seed(42)
```

In [35]: *# Configure paths for the Kaggle environment*

```
DATA_PATH = "/kaggle/input/musicgenreclassification/genres_original"
```

```

# Create genre mapping from folder names
genre_folders = [d for d in os.listdir(DATA_PATH) if os.path.isdir(os.path.join(DATA_PATH, d))]
genre_mapping = {genre: idx for idx, genre in enumerate(sorted(genre_folders))}
num_genres = len(genre_mapping)
print(f"Found {num_genres} genres: {genre_mapping}")

```

Found 10 genres: {'blues': 0, 'classical': 1, 'country': 2, 'disco': 3, 'hiphop': 4, 'jazz': 5, 'metal': 6, 'pop': 7, 'reggae': 8, 'rock': 9}

```

In [36]: # Check the problematic file
problematic_file = "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav"

# Try to load and examine the file
try:
    waveform, sr = torchaudio.load(problematic_file)
    print(f"File loaded successfully:")
    print(f"Sample rate: {sr}")
    print(f"Waveform shape: {waveform.shape}")
    print(f"Min value: {waveform.min()}")
    print(f"Max value: {waveform.max()}")

    # Plot the waveform
    plt.figure(figsize=(12, 4))
    plt.plot(waveform[0])
    plt.title("Waveform of jazz.00054.wav")
    plt.xlabel("Sample")
    plt.ylabel("Amplitude")
    plt.show()

except Exception as e:
    print(f"Error loading file: {e}")
    print("This confirms the file is indeed corrupted")

```

Error loading file: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).
Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):
frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)
frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)
frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)
frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string>>> const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)
frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)
frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)
frame #6: /usr/bin/python3() [0x55559b]
frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)
frame #8: /usr/bin/python3() [0x58536d]
frame #9: /usr/bin/python3() [0x56e229]
frame #10: /usr/bin/python3() [0x52fa60]
frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)
frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)
frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)
frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #15: /usr/bin/python3() [0x56deb6]
frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)
frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)
frame #18: /usr/bin/python3() [0x6135e4]
frame #19: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #20: /usr/bin/python3() [0x62ca33]
frame #21: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #22: /usr/bin/python3() [0x6284b0]
frame #23: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #24: /usr/bin/python3() [0x6284b0]
frame #25: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #26: /usr/bin/python3() [0x6284b0]
frame #27: /usr/bin/python3() [0x62aaec]
frame #28: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)

```
frame #29: /usr/bin/python3() [0x585a87]
frame #30: /usr/bin/python3() [0x58526e]
frame #31: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x6284b0]
frame #34: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x6284b0]
frame #36: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6284b0]
frame #38: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x6284b0]
frame #40: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #43: /usr/bin/python3() [0x553a1f]
frame #44: /usr/bin/python3() [0x4d0bc0]
frame #45: /usr/bin/python3() [0x4e94f3]
frame #46: /usr/bin/python3() [0x54b25b]
frame #47: _PyEval_EvalFrameDefault + 0x9129 (0x546269 in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x6135e4]
frame #49: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #50: /usr/bin/python3() [0x62ca33]
frame #51: /usr/bin/python3() [0x54b25b]
frame #52: PyObject_Vectorcall + 0x35 (0x54b145 in /usr/bin/python3)
frame #53: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)
frame #54: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #55: /usr/bin/python3() [0x63e860]
frame #56: Py_RunMain + 0x13c (0x63e1bc in /usr/bin/python3)
frame #57: Py_BytesMain + 0x2d (0x603f2d in /usr/bin/python3)
frame #58: <unknown function> + 0x29d90 (0x7a5347334d90 in /lib/x86_64-linux-gnu/libc.so.6)
frame #59: __libc_start_main + 0x80 (0x7a5347334e40 in /lib/x86_64-linux-gnu/libc.so.6)
frame #60: _start + 0x25 (0x603db5 in /usr/bin/python3)
```

This confirms the file is indeed corrupted

```
In [37]: # Count files by genre
genre_counts = {}
problem_files = []

for genre in genre_mapping.keys():
    genre_path = os.path.join(DATA_PATH, genre)
```

```
if os.path.isdir(genre_path):
    files = [f for f in os.listdir(genre_path) if f.endswith('.wav') or f.endswith('.au')]
    genre_counts[genre] = len(files)

    # Try to load each file to find any other problematic ones
    for fname in files:
        file_path = os.path.join(genre_path, fname)
        try:
            waveform, sr = torchaudio.load(file_path)
        except Exception as e:
            problem_files.append((file_path, str(e)))

print("Files per genre:")
for genre, count in genre_counts.items():
    print(f"{genre}: {count}")

print("\nProblematic files found:")
for file_path, error in problem_files:
    print(f"{file_path}: {error}")
```

Files per genre:

blues: 100

classical: 100

country: 100

disco: 100

hiphop: 100

jazz: 100

metal: 100

pop: 100

reggae: 100

rock: 100

Problematic files found:

/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string>>> const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

```
frame #15: /usr/bin/python3() [0x56deb6]
frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)
frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)
frame #18: /usr/bin/python3() [0x6135e4]
frame #19: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #20: /usr/bin/python3() [0x62ca33]
frame #21: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #22: /usr/bin/python3() [0x6284b0]
frame #23: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #24: /usr/bin/python3() [0x6284b0]
frame #25: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #26: /usr/bin/python3() [0x6284b0]
frame #27: /usr/bin/python3() [0x62aaec]
frame #28: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #29: /usr/bin/python3() [0x585a87]
frame #30: /usr/bin/python3() [0x58526e]
frame #31: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x6284b0]
frame #34: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x6284b0]
frame #36: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6284b0]
frame #38: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x6284b0]
frame #40: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #43: /usr/bin/python3() [0x553a1f]
frame #44: /usr/bin/python3() [0x4d0bc0]
frame #45: /usr/bin/python3() [0x4e94f3]
frame #46: /usr/bin/python3() [0x54b25b]
frame #47: _PyEval_EvalFrameDefault + 0x9129 (0x546269 in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x6135e4]
frame #49: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #50: /usr/bin/python3() [0x62ca33]
frame #51: /usr/bin/python3() [0x54b25b]
frame #52: PyObject_Vectorcall + 0x35 (0x54b145 in /usr/bin/python3)
frame #53: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)
frame #54: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #55: /usr/bin/python3() [0x63e860]
```

```
frame #56: Py_RunMain + 0x13c (0x63e1bc in /usr/bin/python3)
frame #57: Py_BytesMain + 0x2d (0x603f2d in /usr/bin/python3)
frame #58: <unknown function> + 0x29d90 (0x7a5347334d90 in /lib/x86_64-linux-gnu/libc.so.6)
frame #59: __libc_start_main + 0x80 (0x7a5347334e40 in /lib/x86_64-linux-gnu/libc.so.6)
frame #60: _start + 0x25 (0x603db5 in /usr/bin/python3)
```

In [38]: *# Audio Dataset Class*

```
class AudioGenreDataset(Dataset):
    def __init__(self, audio_files, labels, sample_rate=16000, duration=3.0, transform=None):
        self.audio_files = audio_files
        self.labels = labels
        self.sample_rate = sample_rate
        self.max_length = int(duration * sample_rate) # Convert duration to samples
        self.transform = transform

    def __len__(self):
        return len(self.audio_files)

    def __getitem__(self, idx):
        # Load audio file
        try:
            waveform, sr = torchaudio.load(self.audio_files[idx])

            # Convert to mono if stereo
            if waveform.shape[0] > 1:
                waveform = torch.mean(waveform, dim=0, keepdim=True)

            # Resample if necessary
            if sr != self.sample_rate:
                resampler = torchaudio.transforms.Resample(sr, self.sample_rate)
                waveform = resampler(waveform)

            # Ensure consistent length
            if waveform.shape[1] < self.max_length:
                # Pad if too short
                padding = torch.zeros(1, self.max_length - waveform.shape[1])
                waveform = torch.cat([waveform, padding], dim=1)
            else:
                # Random crop if too long for training variety
                start = torch.randint(0, waveform.shape[1] - self.max_length + 1, (1,))
                waveform = waveform[:, start:start + self.max_length]
```



```

        # Apply transforms if provided
        if self.transform:
            waveform = self.transform(waveform)

        return waveform.squeeze(0), self.labels[idx]

    except Exception as e:
        print(f"Error loading {self.audio_files[idx]}: {e}")
        # Return zeros as a fallback
        return torch.zeros(self.max_length), self.labels[idx]

```

```

In [39]: # Model Architecture using a pretrained model for raw waveform processing
class WaveformGenreClassifier(nn.Module):
    def __init__(self, pretrained_model_name="facebook/wav2vec2-base", num_genres=10, freeze_feature_extractor=False):
        super().__init__()

        # Load the pretrained model
        try:
            self.encoder = Wav2Vec2Model.from_pretrained(pretrained_model_name)
        except Exception as e:
            print(f"Error loading pretrained model: {e}")
            print("Initializing from config instead...")
            config = Wav2Vec2Config.from_pretrained(pretrained_model_name)
            self.encoder = Wav2Vec2Model(config)

        # Freeze the feature extractor if requested
        if freeze_feature_extractor:
            for param in self.encoder.parameters():
                param.requires_grad = False

        # Get the dimension of the encoder's output
        encoder_dim = self.encoder.config.hidden_size # 768 for wav2vec2-base

        # Classification head
        self.classifier = nn.Sequential(
            nn.Linear(encoder_dim, 256),
            nn.BatchNorm1d(256),
            nn.ReLU(),
            nn.Dropout(0.3),
            nn.Linear(256, 128),
            nn.BatchNorm1d(128),

```

```

        nn.ReLU(),
        nn.Dropout(0.2),
        nn.Linear(128, num_genres)
    )

    def forward(self, x):
        # Wav2Vec2 expects inputs of shape [batch_size, sequence_length]
        # Extract features from the encoder
        features = self.encoder(x).last_hidden_state # [batch_size, sequence_length, hidden_size]

        # Global pooling over the sequence length
        pooled = torch.mean(features, dim=1) # [batch_size, hidden_size]

        # Classification
        output = self.classifier(pooled)
        return output

```

```

In [40]: # Data preparation function
def prepare_dataset():
    all_files = []
    all_labels = []

    # Walk through the dataset directory
    for genre in genre_mapping.keys():
        genre_path = os.path.join(DATA_PATH, genre)
        if not os.path.isdir(genre_path):
            continue

        # Get all audio files for this genre
        for file in os.listdir(genre_path):
            if file.endswith('.wav') or file.endswith('.au'):
                file_path = os.path.join(genre_path, file)
                all_files.append(file_path)
                all_labels.append(genre_mapping[genre])

    # Split the dataset: 80% training, 20% testing (no validation set)
    train_files, test_files, train_labels, test_labels = train_test_split(
        all_files, all_labels, test_size=0.2, random_state=42, stratify=all_labels
    )

    # Create datasets
    train_dataset = AudioGenreDataset(train_files, train_labels)

```

```

test_dataset = AudioGenreDataset(test_files, test_labels)

# Create dataloaders
train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True, num_workers=2)
test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False, num_workers=2)

print(f"Train set: {len(train_dataset)} samples")
print(f"Test set: {len(test_dataset)} samples")

return train_loader, test_loader

```

```

In [41]: # Training function
def train_model(model, train_loader, epochs=10, lr=1e-4, weight_decay=1e-5, unfreeze_after=3):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    print(f"Using device: {device}")

    model = model.to(device)
    criterion = nn.CrossEntropyLoss()

    # Initial optimizer with only classification head parameters
    optimizer = optim.AdamW(
        [p for p in model.classifier.parameters() if p.requires_grad],
        lr=lr, weight_decay=weight_decay
    )

    # Learning rate scheduler - use loss from training instead of validation
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(
        optimizer, mode='min', factor=0.5, patience=2, verbose=True
    )

    history = {
        'train_loss': [],
        'train_acc': []
    }

    best_train_loss = float('inf')
    best_model_state = None

    for epoch in range(epochs):
        print(f"\nEpoch {epoch+1}/{epochs}")

        # Unfreeze encoder after specified epochs

```

```

if epoch == unfreeze_after:
    print("Unfreezing encoder layers...")
    for param in model.encoder.parameters():
        param.requires_grad = True

    # Update optimizer to include all parameters
    optimizer = optim.AdamW(
        model.parameters(),
        lr=lr/10, # Lower learning rate for fine-tuning
        weight_decay=weight_decay
    )
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(
        optimizer, mode='min', factor=0.5, patience=2, verbose=True
    )

# Training phase
model.train()
train_loss = 0.0
correct = 0
total = 0

for inputs, labels in tqdm(train_loader, desc="Training"):
    inputs, labels = inputs.to(device), labels.to(device)

    optimizer.zero_grad()

    outputs = model(inputs)
    loss = criterion(outputs, labels)

    loss.backward()
    # Gradient clipping to prevent exploding gradients
    torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    optimizer.step()

    train_loss += loss.item()

    _, predicted = outputs.max(1)
    total += labels.size(0)
    correct += predicted.eq(labels).sum().item()

train_loss = train_loss / len(train_loader)
train_acc = correct / total

```

```

# Update scheduler using training loss (instead of validation loss)
scheduler.step(train_loss)

# Save best model based on training loss
if train_loss < best_train_loss:
    best_train_loss = train_loss
    best_model_state = model.state_dict().copy()
    print(f"New best model saved with train_loss: {train_loss:.4f}")

# Update history
history['train_loss'].append(train_loss)
history['train_acc'].append(train_acc)

print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.4f}")

# Load best model
model.load_state_dict(best_model_state)

# Plot training history
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history['train_loss'], label='Train Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history['train_acc'], label='Train Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training Accuracy')
plt.legend()

plt.tight_layout()
plt.savefig('training_history.png')
plt.show()

return model, history

```

```

In [42]: # Evaluate the model on the test set
def evaluate_model(model, test_loader):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model = model.to(device)
    model.eval()

    all_preds = []
    all_labels = []

    with torch.no_grad():
        for inputs, labels in tqdm(test_loader, desc="Testing"):
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            _, predicted = outputs.max(1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

    # Compute confusion matrix
    cm = confusion_matrix(all_labels, all_preds)

    # Plot confusion matrix
    plt.figure(figsize=(10, 8))
    reverse_mapping = {v: k for k, v in genre_mapping.items()}
    labels = [reverse_mapping[i] for i in range(num_genres)]

    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title('Confusion Matrix')
    plt.tight_layout()
    plt.savefig('confusion_matrix.png')
    plt.show()

    # Print classification report
    report = classification_report(
        all_labels, all_preds,
        target_names=[reverse_mapping[i] for i in range(num_genres)],
        digits=4
    )

```

```

print(report)

# Calculate overall accuracy
accuracy = (np.array(all_preds) == np.array(all_labels)).mean()
print(f"Overall accuracy: {accuracy:.4f}")

return accuracy, cm, report

```

```

In [43]: # Main execution
def main():
    print("Preparing datasets...")
    train_loader, test_loader = prepare_dataset()

    print("Initializing model...")
    model = WaveformGenreClassifier(num_genres=num_genres, freeze_feature_extractor=True)

    print("Training model...")
    trained_model, history = train_model(
        model,
        train_loader,
        epochs=15,
        lr=1e-4,
        weight_decay=1e-5,
        unfreeze_after=5
    )

    # Save the trained model
    torch.save(trained_model.state_dict(), 'genre_classifier_waveform.pth')
    print("Model saved as 'genre_classifier_waveform.pth'")

    print("Evaluating model...")
    accuracy, cm, report = evaluate_model(trained_model, test_loader)

    # Save the evaluation results
    with open('evaluation_results.txt', 'w') as f:
        f.write(f"Accuracy: {accuracy:.4f}\n\n")
        f.write("Classification Report:\n")
        f.write(report)

    print("Done!")

```

```
# Run the main function if executing as script
if __name__ == "__main__":
    main()
```

```
Preparing datasets...
```

Train set: 800 samples

Test set: 200 samples

```
Initializing model...
```

```
/usr/local/lib/python3.11/dist-packages/transformers/configuration_utils.py:311: UserWarning: Passing `gradient_checkpointing` to a config initialization is deprecated and will be removed in v5 Transformers. Using `model.gradient_checkpointing_enable()` instead, or if you are using the `Trainer` API, pass `gradient_checkpointing=True` in your `TrainingArguments`.
```

```
warnings.warn(
```

Training model...

```
Using device: cuda
```

Epoch 1/15

```
Training:  0%|          | 0/50 [00:00<?, ?it/s]
```


Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 2.2524

Train Loss: 2.2524, Train Acc: 0.1525

Epoch 2/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.9651

Train Loss: 1.9651, Train Acc: 0.3337

Epoch 3/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string>>> const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.8161

Train Loss: 1.8161, Train Acc: 0.3937

Epoch 4/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.7080

Train Loss: 1.7080, Train Acc: 0.4512

Epoch 5/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string>>> const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.6597

Train Loss: 1.6597, Train Acc: 0.4763

Epoch 6/15

Unfreezing encoder layers...

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.5941

Train Loss: 1.5941, Train Acc: 0.4888

Epoch 7/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.4911

Train Loss: 1.4911, Train Acc: 0.5437

Epoch 8/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)


```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.4535

Train Loss: 1.4535, Train Acc: 0.5563

Epoch 9/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string>>> const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.4104

Train Loss: 1.4104, Train Acc: 0.5800

Epoch 10/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string>>> const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.3337

Train Loss: 1.3337, Train Acc: 0.6212

Epoch 11/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.3140

Train Loss: 1.3140, Train Acc: 0.6050

Epoch 12/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)


```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.2820

Train Loss: 1.2820, Train Acc: 0.6388

Epoch 13/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.2450

Train Loss: 1.2450, Train Acc: 0.6687

Epoch 14/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

New best model saved with train_loss: 1.1649

Train Loss: 1.1649, Train Acc: 0.7087

Epoch 15/15

Training: 0%| | 0/50 [00:00<?, ?it/s]

Error loading /kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav: Failed to open the input "/kaggle/input/musicgenreclassification/genres_original/jazz/jazz.00054.wav" (Invalid data found when processing input).

Exception raised from get_input_format_context at /__w/audio/audio/pytorch/audio/src/libtorio/ffmpeg/stream_reader/stream_reader.cpp:42 (most recent call first):

frame #0: c10::Error::Error(c10::SourceLocation, std::string) + 0x96 (0x7a52ff4b9446 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #1: c10::detail::torchCheckFail(char const*, char const*, unsigned int, std::string const&) + 0x64 (0x7a52ff4636e4 in /usr/local/lib/python3.11/dist-packages/torch/lib/libc10.so)

frame #2: <unknown function> + 0x42134 (0x7a52b34ca134 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #3: torio::io::StreamingMediaDecoder::StreamingMediaDecoder(std::string const&, std::optional<std::string> const&, std::optional<std::map<std::string, std::string, std::less<std::string>, std::allocator<std::pair<std::string const, std::string> > > const&) + 0x14 (0x7a52b34ccb34 in /usr/local/lib/python3.11/dist-packages/torio/lib/libtorio_ffmpeg4.so)

frame #4: <unknown function> + 0x3a8de (0x7a51fbc488de in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #5: <unknown function> + 0x323ee (0x7a51fbc403ee in /usr/local/lib/python3.11/dist-packages/torio/lib/_torio_ffmpeg4.so)

frame #6: /usr/bin/python3() [0x55559b]

frame #7: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #8: /usr/bin/python3() [0x58536d]

frame #9: /usr/bin/python3() [0x56e229]

frame #10: /usr/bin/python3() [0x52fa60]

frame #11: <unknown function> + 0xfc8b (0x7a5340288c8b in /usr/local/lib/python3.11/dist-packages/torchaudio/lib/_torchaudio.so)

frame #12: _PyObject_MakeTpCall + 0x27c (0x52f67c in /usr/bin/python3)

frame #13: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #14: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #15: /usr/bin/python3() [0x56deb6]

frame #16: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #17: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #18: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #19: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)

frame #20: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

frame #21: /usr/bin/python3() [0x56deb6]

frame #22: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

frame #23: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

frame #24: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)

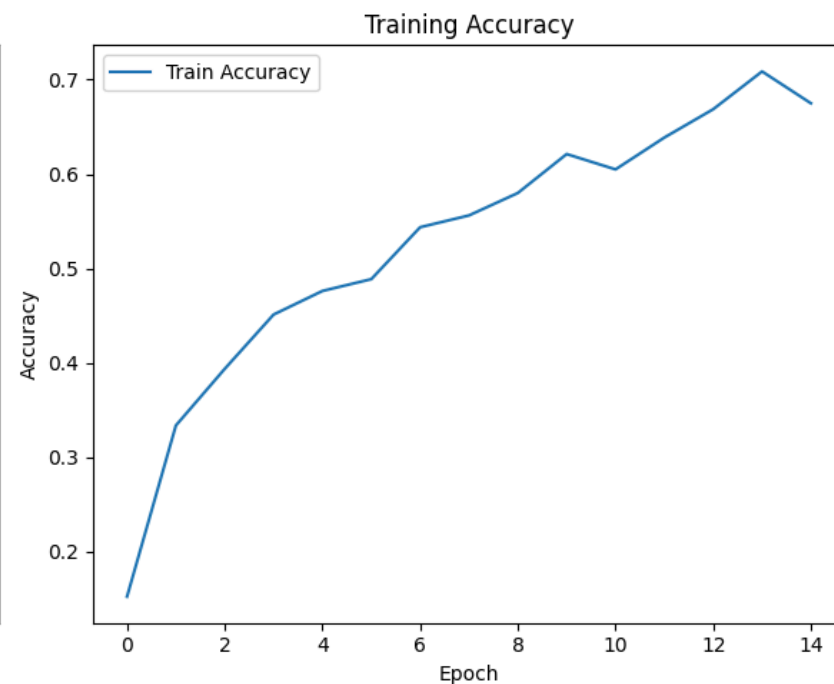
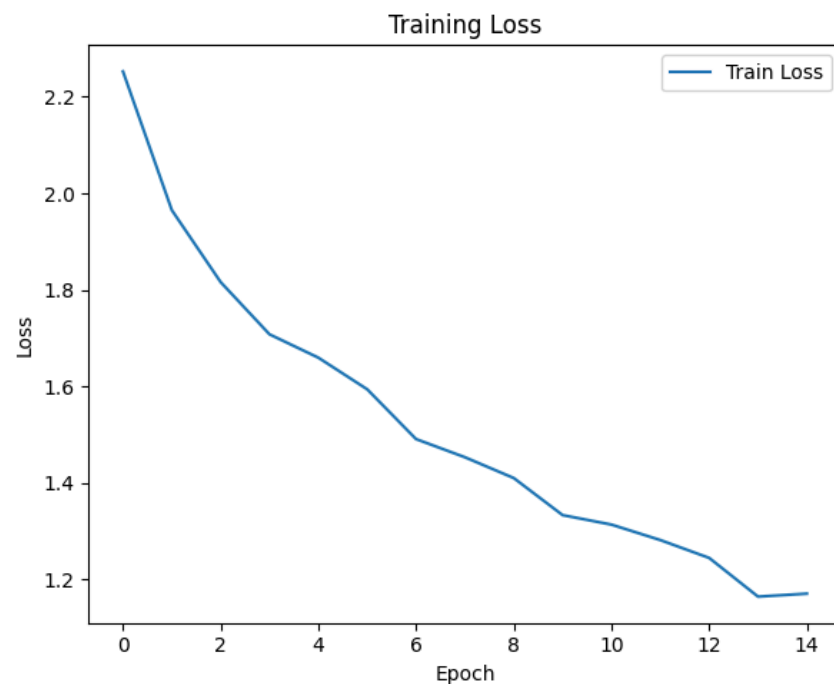
frame #25: /usr/bin/python3() [0x56deb6]

frame #26: _PyObject_MakeTpCall + 0x23b (0x52f63b in /usr/bin/python3)

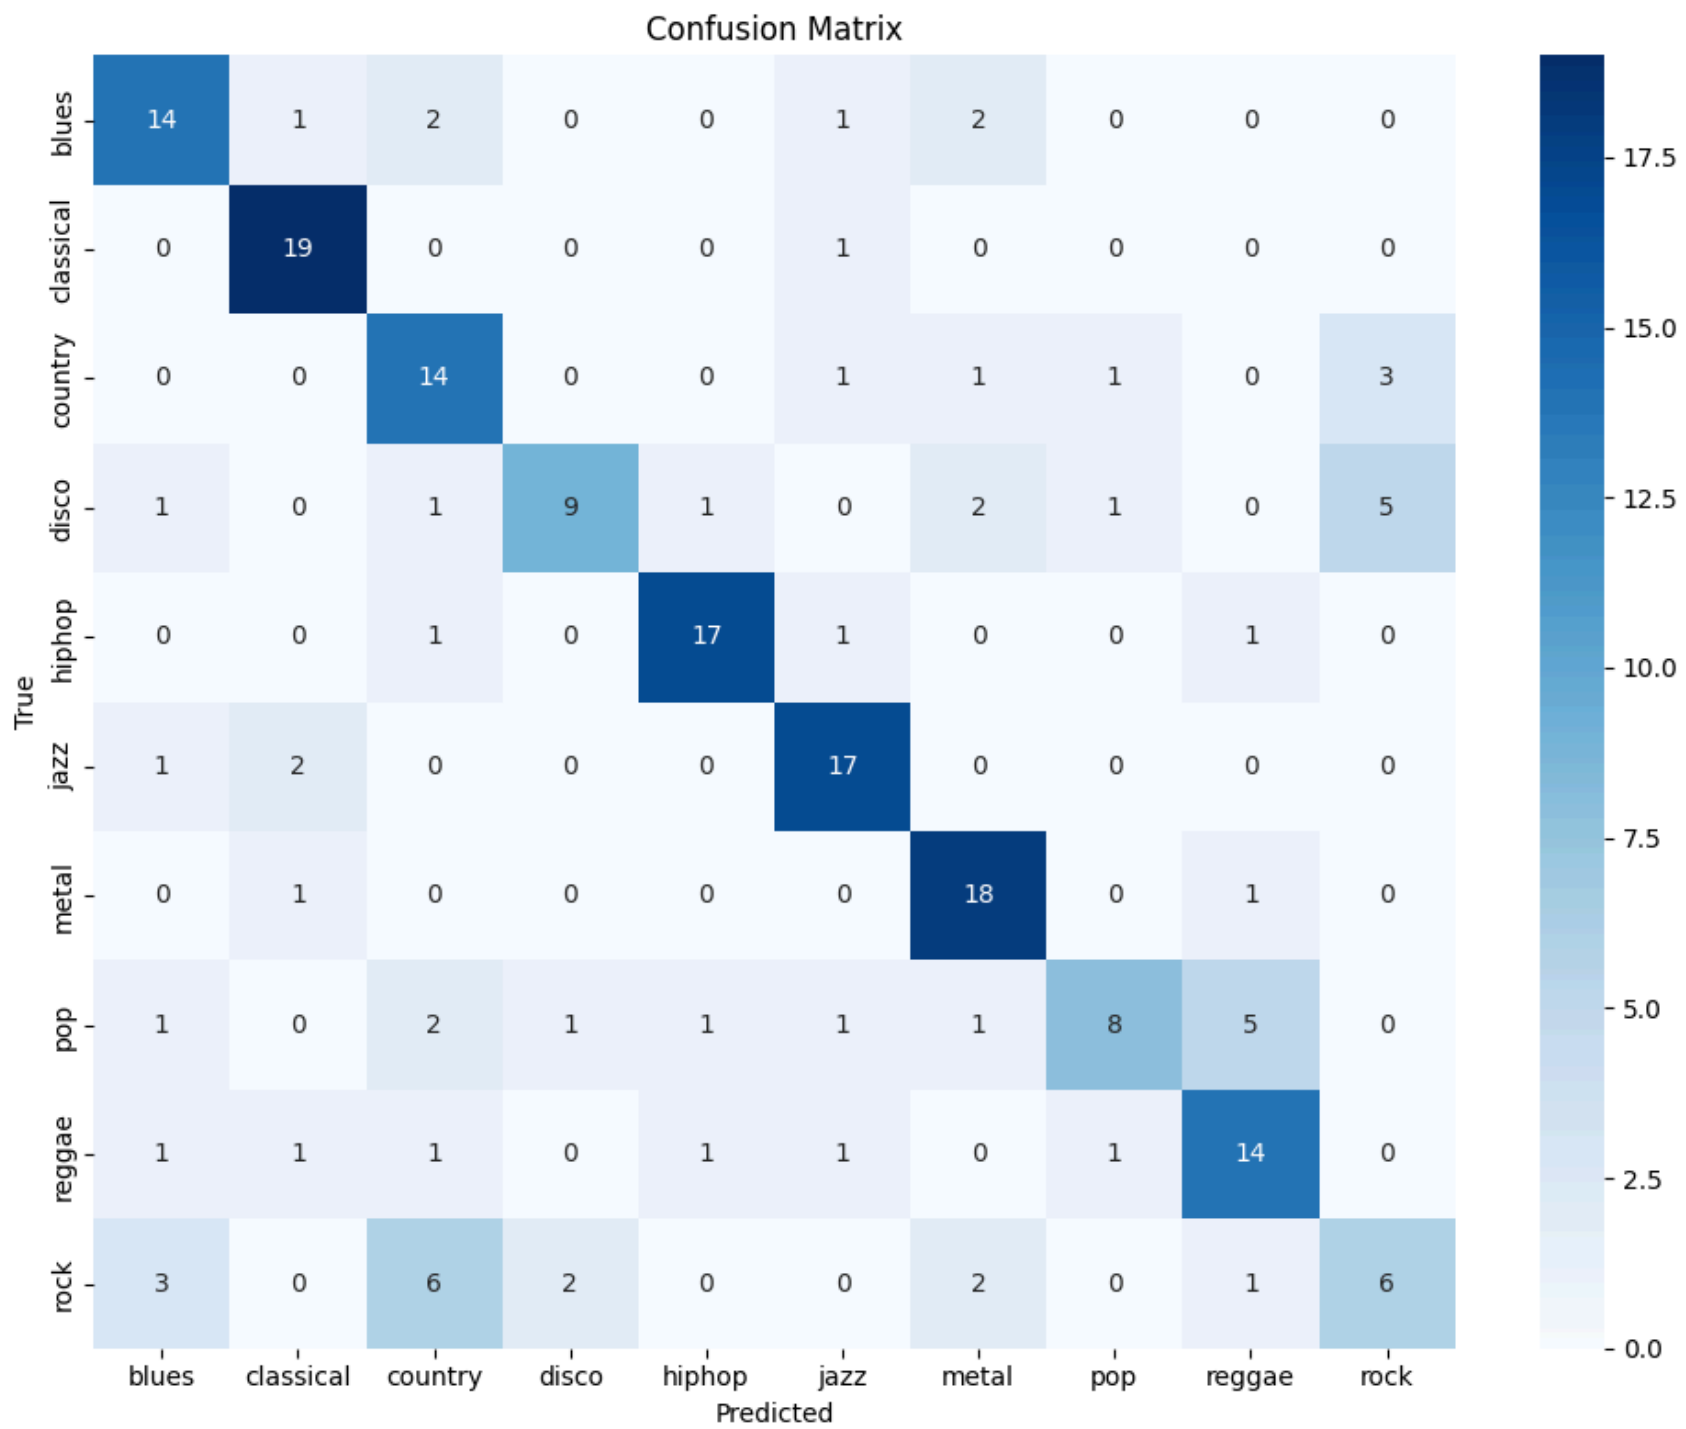
frame #27: _PyEval_EvalFrameDefault + 0x6bf (0x53d7ff in /usr/bin/python3)

```
frame #28: _PyFunction_Vectorcall + 0x173 (0x5661a3 in /usr/bin/python3)
frame #29: PyObject_CallOneArg + 0x47 (0x56bd67 in /usr/bin/python3)
frame #30: /usr/bin/python3() [0x65634d]
frame #31: PyObject_GetIter + 0x18 (0x52c148 in /usr/bin/python3)
frame #32: _PyEval_EvalFrameDefault + 0x2331 (0x53f471 in /usr/bin/python3)
frame #33: /usr/bin/python3() [0x5832c7]
frame #34: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #35: /usr/bin/python3() [0x5832c7]
frame #36: _PyEval_EvalFrameDefault + 0x549 (0x53d689 in /usr/bin/python3)
frame #37: /usr/bin/python3() [0x6135e4]
frame #38: PyEval_EvalCode + 0x97 (0x612c47 in /usr/bin/python3)
frame #39: /usr/bin/python3() [0x62ca33]
frame #40: _PyEval_EvalFrameDefault + 0x390f (0x540a4f in /usr/bin/python3)
frame #41: /usr/bin/python3() [0x6284b0]
frame #42: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #43: /usr/bin/python3() [0x6284b0]
frame #44: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #45: /usr/bin/python3() [0x6284b0]
frame #46: /usr/bin/python3() [0x62aaec]
frame #47: _PyEval_EvalFrameDefault + 0x3a9d (0x540bdd in /usr/bin/python3)
frame #48: /usr/bin/python3() [0x585a87]
frame #49: /usr/bin/python3() [0x58526e]
frame #50: PyObject_Call + 0xf4 (0x570704 in /usr/bin/python3)
frame #51: _PyEval_EvalFrameDefault + 0x4a8f (0x541bcf in /usr/bin/python3)
frame #52: /usr/bin/python3() [0x6284b0]
frame #53: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #54: /usr/bin/python3() [0x6284b0]
frame #55: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #56: /usr/bin/python3() [0x6284b0]
frame #57: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #58: /usr/bin/python3() [0x6284b0]
frame #59: _PyEval_EvalFrameDefault + 0x3485 (0x5405c5 in /usr/bin/python3)
frame #60: /usr/bin/python3() [0x6284b0]
frame #61: <unknown function> + 0x745f (0x7a5345c3545f in /usr/lib/python3.11/lib-dynload/_asyncio.cpython-311-x86_64-linux-gnu.so)
frame #62: /usr/bin/python3() [0x553a1f]
```

Train Loss: 1.1711, Train Acc: 0.6750



Model saved as 'genre_classifier_waveform.pth'
Evaluating model...
Testing: 0%| | 0/13 [00:00<?, ?it/s]



	precision	recall	f1-score	support
blues	0.6667	0.7000	0.6829	20
classical	0.7917	0.9500	0.8636	20
country	0.5185	0.7000	0.5957	20
disco	0.7500	0.4500	0.5625	20
hiphop	0.8500	0.8500	0.8500	20
jazz	0.7391	0.8500	0.7907	20
metal	0.6923	0.9000	0.7826	20
pop	0.7273	0.4000	0.5161	20
reggae	0.6364	0.7000	0.6667	20
rock	0.4286	0.3000	0.3529	20
accuracy			0.6800	200
macro avg	0.6800	0.6800	0.6664	200
weighted avg	0.6800	0.6800	0.6664	200

Overall accuracy: 0.6800

Done!